

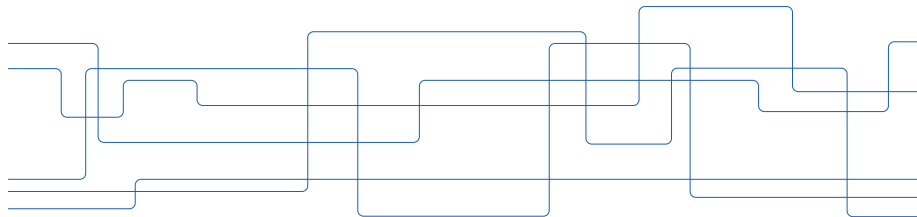


Performance Optimization

Dirk Pleiter

PDC and CST | EECS | KTH

January 2023





Overview

Computational Patterns

Memory Access Optimization

Cache Blocking

(Auto-)Vectorization



Content

Computational Patterns

Memory Access Optimization

Cache Blocking

(Auto-)Vectorization

Introduction to Computational Dwarfs

- ▶ **Dwarf** = An algorithmic method that captures a pattern of computation and communication
 - ▶ Terminology introduced in Krste Asanovic et al., “The Landscape of Parallel Computing Research: A View from Berkeley”, 2006
 - ▶ Various dwarfs already identified in earlier work
- ▶ Goal: Identify commonalities between numerical applications to exploit common knowledge and experience about
 - ▶ Performance characteristics (e.g. arithmetic intensity)
 - ▶ Best-practices for implementation on a given architecture

Computational Dwarfs: Overview

Dwarfs for HPC

1. Dense linear algebra
2. Sparse linear algebra
3. Spectral methods
4. N-body methods
5. Structured grids
6. Unstructured grids
7. MapReduce (Monte Carlo)

Other Dwarfs

8. Combinational logic
9. Graph traversal
10. Dynamic programming
11. Back-track and branch + bound
12. Graphical models
13. Finite state machines

Computational Dwarf: Structured Grids (1/3)

- ▶ Applications perform periodic updates of regular, multi-dimensional grids
- ▶ Memory access features
 - ▶ Regular strided memory access
 - ▶ High spatial data locality, i.e. consecutive access to data that is close in memory
- ▶ Application areas
 - ▶ Heat transfer
 - ▶ Lattice Quantum Chromodynamics
 - ▶ Computational fluid dynamics using the Lattice Boltzmann Method

Computational Dwarf: Structured Grids (2/3)

Example: 2-dimensional Poisson equation

- Formulation in the continuum

$$-\frac{\partial^2 v(x, y)}{\partial x^2} - \frac{\partial^2 v(x, y)}{\partial y^2} = f(x, y)$$

- Discretisation of 2nd-order derivative

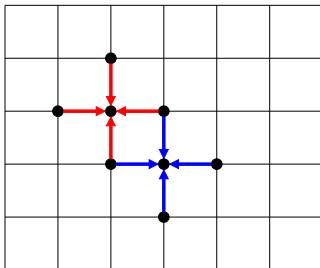
$$-\frac{\partial^2 v(x, y)}{\partial x^2} \leftarrow \frac{2v_{i,j} - v_{i-1,j} - v_{i+1,j}}{h^2}$$

- Discrete Poisson equation in 2 dimensions:


$$T v = h^2 f \quad \text{where} \quad T = \begin{pmatrix} 4 & -1 & 0 & \cdots \\ -1 & 4 & -1 & \cdots \\ \vdots & \vdots & \ddots & \end{pmatrix}$$

Computational Dwarf: Structured Grids (3/3)

Graphical representation of $T v$:



Observations:

- ▶ Matrix T acts as a stencil operator
- ▶ Any element of vector v is reused 4 times  data locality

Computational Dwarf: Dense Linear Algebra

- ▶ Data are densely populated matrices or vectors
$$\begin{pmatrix} 1 & 4 & 5 & 6 \\ 0 & 3 & 2 & 6 \\ 1 & 2 & 3 & 1 \\ 5 & 4 & 3 & 2 \end{pmatrix} \begin{pmatrix} 1 \\ 2 \\ 3 \\ 4 \end{pmatrix} = \begin{pmatrix} 48 \\ 36 \\ 18 \\ 30 \end{pmatrix}$$
- ▶ Typical memory access pattern:
Unit-stride or fixed-stride regular sequential access
- ▶ Example: $y_i \leftarrow \sum_j A_{ij} \cdot x_j$
- ▶ Popular library with dense linear algebra kernels: BLAS (Basic Linear Algebra Subprograms)
 - ▶ Different implementations available:
 - ▶ Open source: NETLIB BLAS, OpenBLAS, BLIS
 - ▶ Closed source: MKL (Intel), ESSL (IBM), Arm Performance Library (Arm)
 - ▶ Classification of operations
 - ▶ BLAS-1: scalar, vector and vector-vector
 - ▶ BLAS-2: matrix-vector
 - ▶ BLAS-3: matrix-matrix



Content

Computational Patterns

Memory Access Optimization

Cache Blocking

(Auto-)Vectorization

Digression: Memory Access Locality

- ▶ Empirical observation: Programs tend to reuse data and instructions they have used recently
- ▶ Observation can be exploited to
 - ▶ Improve performance
 - ▶ Optimize use of more/less expensive memory
- ▶ Types of localities
 - ▶ **Temporal locality**: Recently accessed items are likely to be accessed in the near future
 - ▶ **Spatial locality**: Items whose addresses are near one another tend to be referenced close together in time

Memory Access Locality: Example

```
1 double a[N][N], b[N][N];  
2  
3 for (i=0; i<N; i++)  
4     for (j=1; j<N; j++)  
5         a[i][j] = b[j-1][0] + b[j][0];
```

- ▶ Assume right-most index being fastest running index
- ▶ Temporal locality: $b[j][0]$
- ▶ Spatial locality: $a[i][j]$ and $a[i][j+1]$ ($j+1 < N$)

Memory Access Pattern Optimizations: Stride-1 Memory Access

```
1 double a[N][N], b[N][N];
2
3 for (int j = 0; j < N; j++)
4     for (int i = 0; i < N; i++)
5         b[i][j] = a[i][j] + 0.1;
```

```
1 double a[N][N], b[N][N];
2
3 for (int i = 0; i < N; i++)
4     for (int j = 0; j < N; j++)
5         b[i][j] = a[i][j] + 0.1;
```

- ▶ Assume N to be very large
- ▶ Question: Which of the codes written in C will be faster?
Why?

Stride-1 Memory Access (cont.)

```

1 double a[N][N], b[N][N];
2
3 for (int j = 0; j < N; j++)
4   for (int i = 0; i < N; i++)
5     b[i][j] = a[i][j] + 0.1;


```

```

1 double a[N][N], b[N][N];
2
3 for (int i = 0; i < N; i++)
4   for (int j = 0; j < N; j++)
5     b[i][j] = a[i][j] + 0.1;

```

- ▶ Answer: Right code because of stride-1 memory access
 - ▶ During each inner loop iteration of the left code 2 new cache lines are accessed
 - ▶ Because N is large, cache lines will be evicted before being reused
- ▶ The arrays are mapped to memory as follows assuming a cache line size of 64 Byte:



Memory Access Pattern Optimizations: Indirect Memory Access

Example: Sum over neighbours

```
1 double a[N], b[N];
2 int nnb[N];
3 int nb[N][NNB_MAX];
4
5 for (int i = 0; i < N; i++) {
6     b[i] = 0;
7     for (int j = 0; j < nnb[i]; j++)
8         b[i] += a[nb[i][j]];
9 }
```

- ▶ Performance challenges
 - ▶ Memory latency becomes an issue as `a[]` can only be loaded after the load of `nb[] []` completed
 - ▶ Access to `a[]` likely not sequential with low stride
 - ▶ Vectorization of the code requires gather load instructions and the compiler auto-vectorizer may fail
- ▶ Possible optimizations strategies
 - ▶ Recompute index of neighbour
 - ▶ Improve spatial data locality by sorting data



Content

Computational Patterns

Memory Access Optimization

Cache Blocking



(Auto-)Vectorization

Cache Blocking

- ▶ **Cache blocking** is a strategy where data structures or memory access patterns are changed such that capacity cache misses are reduced
 - ▶ Optimisation strategy sometimes also called **loop tiling** or **loop blocking**
- ▶ Example: Double-precision matrix-vector multiplication (BLAS-2)

$$y_i \leftarrow \sum_{j=0}^{N-1} A_{ij} x_j \quad (i = 0, \dots, N-1)$$

- ▶ Advantage: x needs to be loaded only once per block
- ▶ Disadvantage: Non-linear read of A



Cache-blocked Matrix-Vector Multiplication

```
1 double x[N], y[N];
2 double A[N][N];
3
4 for (int ib = 0; ib < N/B; ib++)
5 {
6     for (int i = 0; i < B; i++)
7         y[ib*B+i] = 0.0;
8
9     for (int jb = 0; jb < N/B; jb++)
10        for (int i = 0; i < B; i++)
11            for (int j = 0; j < B; j++)
12                {
13                    int ii = ib*B + i;
14                    int jj = jb*B + j;
15                    y[ii] += A[ii][jj] * x[jj];
16                }
17 }
```


Cache-blocked Matrix-Vector Multiplication

- ▶ Let B be the block size with $B \leq N$ and N being a multiple of B
- ▶ Information exchange analysis:

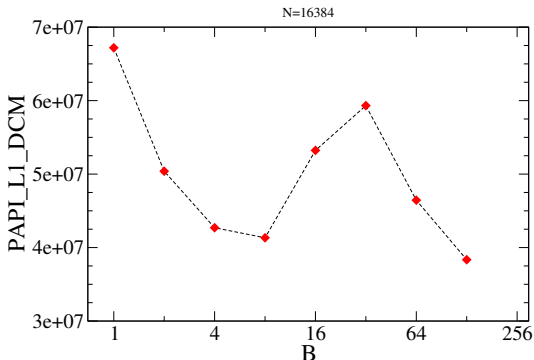
$$\begin{aligned}I_{\text{rf,rf}}(N, B) &= [N + (N - 1)] N \text{ Flop} \\I_{\text{mem,rf}}(N, B) &= (N + N^2 + N \cdot \mathbf{N/B}) 8 \text{ Byte} \\AI &= I_{\text{rf,rf}}/I_{\text{mem,rf}} \simeq \\&\quad (0.125 \dots 0.250) \text{ Flop/Byte}\end{aligned}$$

Cache-blocked Matrix-Vector Multiplication

- ▶ Results obtained for $N = 16384$ on Intel Xeon E5-2623 v3 using a single thread
 - ▶ Horizontal line shows results for no cache blocking
- ▶ Observations
 - ▶ Large number of nested loops generate a significant overhead
 - ▶ Up to $2\times$ speed-up compared to no cache blocking
 - ▶ Best results are obtained for small block sizes with $B = 8$



Cache-blocked Matrix-Vector Multiplication



- ▶ For small B number of L1 data cache misses drop as expected
- ▶ For large B number of L1 data cache misses drop due to memory pre-fetcher



Content

Computational Patterns

Memory Access Optimization

Cache Blocking

(Auto-)Vectorization

SIMD Parallelism

- ▶ Single Instruction Multiple Data (SIMD) instructions exploit data-level parallelism by operating on data items in parallel
 - ▶ E.g., SIMD add

$$\begin{pmatrix} z_0 \\ z_1 \\ z_2 \\ z_3 \end{pmatrix} \leftarrow \begin{pmatrix} x_0 \\ x_1 \\ x_2 \\ x_3 \end{pmatrix} + \begin{pmatrix} y_0 \\ y_1 \\ y_2 \\ y_3 \end{pmatrix}$$

- ▶ Example ISA:
 - ▶ Intel Streaming SIMD Extensions (SSE)
 - ▶ Intel Advanced Vector Extensions (AVX, AVX2, AVX512)
 - ▶ POWER ISA (VMX/Altivec, VSX)
 - ▶ Armv7 NEON, Armv8, Arm SVE

SIMD Programming: Auto-Vectorisation

Auto-vectorisation = Compiler capability to convert operations on scalars to operations on vectors

- ▶ Advantages:
 - ▶ Portable code
 - ▶ Programmer relieved from complex code transformations
- ▶ Disadvantage:
 - ▶ Only indirect control on code generation
 - ▶ Compiler may fail to identify vectorisation opportunities

Auto-Vectorisation: Example (1/3)

- ▶ daxpy (BLAS-1): $\vec{y} \leftarrow \alpha \vec{x} + \vec{y}$

```
1 #define N 1024
2
3 void daxpy(double* x, double alpha, double* y)
4 {
5     for (int i = 0; i < N; i++)
6         y[i] += alpha * x[i];
7 }
```

- ▶ Using GCC compiler in a verbose mode:

```
% gcc -O3 -fopt-info -S daxpy1.c
...
daxpy1.c:6:21: optimized: loop vectorized using 16 byte vectors
daxpy1.c:6:21: optimized: loop versioned for vectorization because of possible
aliasing
daxpy1.c:3:6: note: vectorized 1 loops in function.
daxpy1.c:8:1: note: ***** Analysis failed with vector mode V2DF
daxpy1.c:8:1: note: ***** Skipping vector mode V16QI, which would repeat the
analysis for V2DF
```

- ▶ Compiler generates additional code to check possible aliasing issues at run-time

Auto-Vectorisation: Example (2/3)

- ▶ For daxpy we know that \vec{x} and \vec{y} are stored in different memory locations
- ▶ Use the `restrict` qualifier to inform the compiler about this:

```
1 #define N 1024
2
3 void daxpy(double* restrict x, double alpha, double* restrict y)
4 {
5     for (int i = 0; i < N; i++)
6         y[i] += alpha * x[i];
7 }
```

- ▶ Recompiling modified code:

```
% gcc -O3 -fopt-info -S daxpy1.c
...
daxpy2.c:6:21: optimized: loop vectorized using 16 byte vectors
daxpy2.c:3:6: note: vectorized 1 loops in function.
daxpy2.c:8:1: note: ***** Analysis failed with vector mode VOID
```


Auto-Vectorisation: Example (3/3)

- ▶ Assembler generated for the Kunpeng 920 processor:

```
daxpy:
.LFB0:
    .cfi_startproc
    dup        v2.2d, v0.d[0]
    mov        x2, 0
    .p2align 3,,7
.L2:
    ldr        q0, [x1, x2]
    ldr        q1, [x0, x2]
    fmla       v0.2d, v2.2d, v1.2d
    str        q0, [x1, x2]
    add        x2, x2, 16
    cmp        x2, 8192
    bne        .L2
    ret
    .cfi_endproc
```

- ▶ Advice for reading the assembler:
 - ▶ Code uses 128-bit SIMD registers (suffix .2d)
 - ▶ Address pointer stored in register x2 is incremented by 16 Byte in each loop iteration

Finish with a Simple Architecture: Leibniz' Reckoner



[Museum Schloss Herrenhausen]