

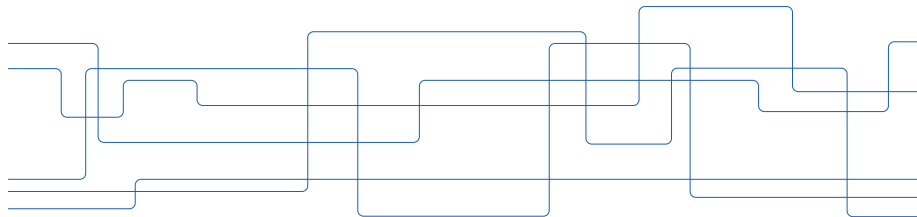


Performance Analysis

Dirk Pleiter

PDC and CST | EECS | KTH

January 2023





Overview

Introduction to Performance

Performance Measurement



Content


Introduction to Performance

Performance Measurement

- ▶ Reliable measure of **performance** = $\frac{\text{amount of work}}{\text{execution time}}$
- ▶ Execution time = time to execute a particular amount of work
- ▶ Ambigiouties in definition of execution time:
 - ▶ Wall-clock time = latency to complete task
 - ▶ Resource occupation time, e.g. CPU time

Bandwidth or Throughput vs. Latency

- ▶ Definition **bandwidth or throughput**
 - ▶ Bandwidth = amount of items passing a particular interface per time unit
 - ▶ Throughput = amount of work done in a given time
- ▶ Definition **latency (or response time) Δt** =
Time between events indicating start of a task/operation and the event signaling its completion
- ▶ Example: Memory read operation




- ▶ Bandwidth: Amount of data passing memory interface per time unit
- ▶ Latency: Time from starting memory read operation until last data item arrived in register file

Latency-Bandwidth Model

- ▶ Simple model describing **latency** Δt as a function of the information exchange I
- ▶ Model parameters:
 - ▶ **Start-up latency** λ = Time between event trigger and start of arrival of response
 - ▶ **Bandwidth** β : Parametrizes time to exchange information I
- ▶ Model:



$$\Delta t(I) = \lambda + \frac{I}{\beta}$$



Latency-Bandwidth Model (2)

- ▶ **Effective bandwidth** $\tilde{b} = I / \Delta t(I)$
- ▶ Example: Read I Bytes from memory or network

$$\Delta t(I) = \lambda + \frac{I}{\beta} \Rightarrow \tilde{b}(I) = \left(\frac{\lambda}{I} + \frac{1}{\beta} \right)^{-1}$$



Information Exchange Function

- ▶ A computation implies that information is transferred from a storage device x to a storage device y .
- ▶ **Information Exchange Function:**

$I_{x,y}^k(W)$ = data transferred between computer sub-systems for specific computation k


x ... source storage device (e.g. memory)

y ... destination storage device (e.g. register file)

W ... **problem size/work-load**

Information Exchange Model

- ▶ **Machine** = Set of interconnected devices
 - ▶ Storage devices
 - ▶ Processing/transport devices
- ▶ **Storage devices**
 - ▶ Examples: Memory, register file, cache
 - ▶ Parameters: Storage size σ
- ▶ **Processing/transport devices**
 - ▶ Examples: Arithmetic pipeline, bus
 - ▶ Parameters: bandwidth/throughput β , startup latency λ
- ▶ **Graphical representation**
 - ▶ Vertices = Storage devices
 - ▶ Edges = Processing/transport devices



Information Exchange Model: Latency Predictions

Ansatz to predict latency:

$$\Delta t_{x,y}^k \simeq \lambda_{x,y} + l_{x,y}^k / \beta_{x,y}$$

Example:

- ▶ $x, y = R$
- ▶ $\beta_{R,R}$ = throughput arithmetic unit
- ▶ $l_{x,y}^k$ = number of input (or output) operands



arithmetic unit

register file

Beware of limitations of this ansatz:

- ▶ Transfer mechanism may depend on task size N
- ▶ Bandwidth changes due to resource congestion is ignored
- ▶ ...

Information Exchange Model: Simple Example (1/2)

- ▶ Example operation: $y_i \leftarrow \alpha \cdot x_i$ ($i = 1, \dots, N$)
 $x_i, y_i \dots$ vectors of single precision floating-point numbers
 $\alpha \dots$ single-precision floating-point number
- ▶ Information exchange:

load x , store y	$I_{\text{mem}} = N \cdot (4 + 4) \text{ Bytes}$
multiply α and x	$I_{\text{fp}} = N \cdot 1 \text{ Flop}$

- ▶ Assume the following hardware parameters:

Memory bandwidth	$\beta_{\text{mem}} = 1 \text{ Byte/clock cycle}$
Floating-point unit throughput	$\beta_{\text{fp}} = 1 \text{ Flop/clock cycle}$

- ▶ Latency predictions (ignoring start-up latency):

Δt_{mem}	$N \cdot 8 \text{ clock cycles}$
Δt_{fp}	$N \cdot 1 \text{ clock cycles}$

Information Exchange Model: Simple Example (2/2)

Latency for full operation:

- ▶ No overlap of memory load/store and arithmetic operations:

$$\Delta t(N) = \Delta t_{\text{mem}} + \Delta t_{\text{fp}} = N \cdot 9 \text{ cycle}$$

- ▶ Perfect overlap of memory load/store and arithmetic operations:

$$\Delta t(N) = \max(\Delta t_{\text{mem}}, \Delta t_{\text{fp}}) = N \cdot 8 \text{ cycle}$$

👉 **Memory bandwidth limited problem**

Memory Bandwidth

▶ Definitions

- ▶ Measured bandwidth b_{mem}
 - ▶ Amount of transferred data I_{mem} over measured latency Δt
- ▶ Latency-bandwidth model parameter β_{mem}
- ▶ Effective bandwidth \tilde{b}_{mem}
- ▶ **Nominal hardware bandwidth** B_{mem}
 - ▶ Data rate on hardware link/bus

▶ Relations

- ▶ Upper limit for measured bandwidth: $b_{\text{mem}} \leq B_{\text{mem}}$
- ▶ Upper limit for effective bandwidth: $\tilde{b}_{\text{mem}} \leq \beta_{\text{mem}}$
- ▶ Note: $\beta_{\text{mem}} \leq B_{\text{mem}}$ is not guaranteed
 - ▶ But expected if bandwidth-latency model is a good model

Arithmetic Intensity

► **Arithmetic Intensity** =

[H. Harris, 2005]

$$AI = \frac{\text{Number of floating-point operations}}{\text{Amount of transferred data}} = \frac{l_{\text{fp}}}{l_{\text{mem}}}$$

► Example: $y_i \leftarrow \alpha \cdot x_i$

$$AI = \frac{1 \text{ Flop}}{8 \text{ Bytes}} \quad \text{single-precision}$$

$$AI = \frac{1 \text{ Flop}}{16 \text{ Bytes}} \quad \text{double-precision}$$

Roofline Model (1/2)

[S. Williams et al., 2009]

- ▶ Floating-point and memory performance limits:

$$b_{\text{fp}} \leq B_{\text{fp}}, \quad b_{\text{mem}} \leq B_{\text{mem}}$$

- ▶ Upper limit for latency assuming perfect overlap of memory and arithmetic operations assuming the latency-bandwidth model to hold with start-up latencies:


$$\Delta t = \max \left(\frac{l_{\text{fp}}}{b_{\text{fp}}}, \frac{l_{\text{mem}}}{b_{\text{mem}}} \right) \geq \max \left(\frac{l_{\text{fp}}}{B_{\text{fp}}}, \frac{l_{\text{mem}}}{B_{\text{mem}}} \right)$$

- ▶ Upper limit (=roof) for floating-point performance (“attainable performance”)

$$b_{\text{fp}} = \frac{l_{\text{fp}}}{\Delta t} \leq \min \left(B_{\text{fp}}, \frac{l_{\text{fp}}}{l_{\text{mem}}} B_{\text{mem}} \right) = \min (B_{\text{fp}}, \text{AI} \cdot B_{\text{mem}})$$

Roofline Model (2/2)

Example: $y_i \leftarrow \alpha \cdot x_i$





Content

Introduction to Performance

Performance Measurement

- ▶ Reliable measure of **performance** = $\frac{\text{amount of work}}{\text{execution time}}$
- ▶ Execution time = time to execute a particular amount of work
- ▶ Ambigiouties in definition of execution time:
 - ▶ Wall-clock time = latency to complete task
 - ▶ Resource occupation time, e.g. CPU time

Performance measurement

- ▶ Types
 - ▶ **Profiling**

Characterize behaviour of an application in terms of aggregate performance metrics

 - ▶ Example: Average time spent in subroutine
 - ▶ **Sampling**

Creation of profiles by means of period or random probing of performance numbers independently of the current state of application execution.
 - ▶ **Event-tracing**

Generate list of event records comprising of time stamp, location, type information, event-specific information
- ▶ Typically additional instructions need to be inserted into program (“instrumentation”)

Performance Measurement (2)

▶ Hardware support for performance measurement

- ▶ Modern processors provide hardware counters organised in a Performance Monitoring Unit (PMU)

Hardware counters = Small set of registers that count events

- ▶ Typical events
 - ▶ Clock tick
 - ▶ Cache miss
 - ▶ Instruction scheduled

▶ Issues related to performance measurement

- ▶ Code instrumentation may impact performance
- ▶ Accessibility of hardware counters
 - ▶ Typically support from kernel is required
 - ▶ System-wide services may be used for collecting data
- ▶ Management of large event traces

Typical Time Scales

Example: Intel Nehalem architecture

Core clock cycle	200 – 300 ps
L1 cache access latency	~ 5 cycles
L2 cache access latency	~ 10 cycles
L3 cache access latency	$\sim 30 - 50$ cycles
Memory access latency	200 – 300 cycles

☞ Need time resolution $\ll 1 \mu\text{sec}$

Time Measurement

- ▶ `/usr/bin/time`
 - ▶ Wall, user, system time
 - ▶ Resolution given by ticks per second:
`sysconf(_SC_CLK_TCK)` typically: 10 ms
- ▶ `clock_gettime(clockid_t, struct timespec *)`
 - ▶ Support for different clocks, recommended to use `CLOCK_PROCESS_CPUTIME_ID` or `CLOCK_THREAD_CPUTIME_ID`
 - ▶ Returns current clock as structure

```
struct timespec {  
    time_t tv_sec;    /* seconds */  
    long tv_nsec;    /* nanoseconds */  
};
```
 - ▶ Include file `time.h` and link with `-lrt`

Performance Tools: PAPI

PAPI = Performance Application Programming Interface

- ▶ Developed at University of Tennessee
- ▶ De facto standard to interface with PMU on various architectures
 - ▶ Examples: Intel Xeon, Xeon Phi; AMD; IBM POWER; ARM
- ▶ Selected counters supported by PAPI:

PAPI_TOT_CYC	Total cycles
PAPI_TOT_INS	Instructions completed
PAPI_FP_INS	Floating point instructions
PAPI_L1_DCM	Level 1 data cache misses
PAPI_L2_DCM	Level 2 data cache misses
PAPI_L3_DCM	Level 3 data cache misses
PAPI_BR_INS	Branch instructions
PAPI_BR_MSP	Conditional branch instructions mispredicted

PAPI: Example Code

```
#include <stdlib.h>
#include <stdio.h>
#include <papi.h>

# define N 10000
double a[N], b[N];

void kernel() {
    for (int i = 0; i < N; i++)
        b[i] = a[i] + 0.1;
}

int main() {
    // Start measurement
    if (PAPI_hl_region_begin(" kernel") != PAPI_OK)
        exit(1);

    // Some kernel
    kernel();

    // Read counters
    if (PAPI_hl_region_end(" kernel") != PAPI_OK)
        exit(1);

    return 0;
}
```

PAPI: Example Execution

Compile example program:

```
% gcc -o ex.x -lpapi
```

Execute example program:

```
% export PAPI_EVENTS="PAPI_TOT_CYC,PAPI_LD_INS"  
% export PAPI_OUTPUT_DIRECTORY="."  
% ./ex.x
```

```
{  
  "papi_version":"7.0.0.0",  
  "cpu_info":" Hisilicon Kunpeng",  
  "max_cpu_rate_mhz":"2600",  
  "min_cpu_rate_mhz":"200",  
  "event_definitions":{  
    "PAPI_TOT_CYC":{  
      "component":" perf_event",  
      "type":" delta"  
    },  
    "PAPI_LD_INS":{  
      "component":" perf_event",  
      "type":" delta"  
    }  
  },  
  "threads":{  
    "0":{  
      "regions":{  
        "0":{  
          "name":" kernel",  
          "parent_region_id":" -1",  
          "cycles":" 5293",  
          "real_time_nsec":" 50140",  
          "PAPI_TOT_CYC":" 106057",  
          "PAPI_LD_INS":" 61411"  
        }  
      }  
    }  
  }  
}
```

PAPI: Example Performance Analysis

- ▶ For $N = 10000$ loop iterations we measured $N_{ld} = 61411$
 $\Rightarrow n_{ld} = N_{ld}/N \simeq 6$
- ▶ Inspection of the assembler code: Number of load instructions per loop iteration is $n_{ld} = 6$
 - ▶ Tip: Use `gcc -S ex.c` to obtain the assembler code
- ▶ Inspection of the C code: Only need to load `a[i]`, i.e. the optimum would be $n_{ld} = 1$

```
kernel:
.LFB6:
    .cfi_startproc
    sub     sp, sp, #16
    .cfi_def_cfa_offset 16
    str     wzr, [sp, 12]
    b       .L2

.L3:
    adrp    x0, a
    add     x0, x0, :lo12:a
    ldrsw   x1, [sp, 12]
    ldr     d0, [x0, x1, lsl #3]
    adrp    x0, .LC0
    ldr     d1, [x0, #:lo12:.LC0]
    fadd    d0, d0, d1
    adrp    x0, b
    add     x0, x0, :lo12:b
    ldrsw   x1, [sp, 12]
    str     d0, [x0, x1, lsl #3]
    ldr     w0, [sp, 12]
    add     w0, w0, 1
    str     w0, [sp, 12]

.L2:
    ldr     w1, [sp, 12]
    mov     w0, 9999
    cmp     w1, w0
    ble     .L3
    nop
    nop
    add     sp, sp, 16
    .cfi_def_cfa_offset 0
    ret
```

PAPI: Example Execution (Optimized)

Compile example program with more optimizations enabled:

```
% gcc -O3 -o ex-03.x -lpapi
```

Execute example program:

```
% export  
PAPI_EVENTS="PAPI_TOT_CYC,PAPI_LD_INS"  
% export PAPI_OUTPUT_DIRECTORY="."  
% ./ex-03.x
```

```
{  
  "papi_version": "7.0.0.0",  
  "cpu_info": "Hisilicon Kunpeng",  
  "max_cpu_rate_mhz": "2600",  
  "min_cpu_rate_mhz": "200",  
  "event_definitions": {  
    "PAPI_TOT_CYC": {  
      "component": "perf_event",  
      "type": "delta"  
    },  
    "PAPI_LD_INS": {  
      "component": "perf_event",  
      "type": "delta"  
    }  
  },  
  "threads": {  
    "0": {  
      "regions": {  
        "0": {  
          "name": "kernel",  
          "parent_region_id": "-1",  
          "cycles": "2500",  
          "real_time_nsec": "22840",  
          "PAPI_TOT_CYC": "40591",  
          "PAPI_LD_INS": "6297"  
        }  
      }  
    }  
  }  
}
```

PAPI: Example Performance Analysis (Optimized)

- ▶ For $N = 10000$ loop iterations we measured $N_{ld} = 6297$
 $\Rightarrow n_{ld} = N_{ld}/N \simeq 0.5$
- ▶ Inspection of the assembler code: Number of load instructions per loop iteration is $n_{ld} = 1$
- ▶ Question: Why do we observe less load instructions than expected?

```
kernel:
.LFB22:
    .cfi_startproc
    adrp    x0, .LC0
    adrp    x2, b
    adrp    x1, a
    mov     x3, 14464
    add     x2, x2, :lo12:b
    add     x1, x1, :lo12:a
    ldr     q1, [x0, #:lo12:.LC0]
    movk    x3, 0x1, lsl 16
    mov     x0, 0
    .p2align 3,,7
.L2:
    ldr     q0, [x1, x0]
    fadd    v0.2d, v0.2d, v1.2d
    str     q0, [x2, x0]
    add     x0, x0, 16
    cmp     x0, x3
    bne     .L2
    ret
```

Performance Tools: Perf

perf = Performance analysis tool for Linux

- Uses the Performance Counters for Linux (PCL)
- Listing of available events: `perf list`
- Example usage:

```
% perf stat ./ex-03.x
```

```
Performance counter stats for './ex-03.x':
```

```


64.14 msec task-clock:u
0 context-switches:u
0 cpu-migrations:u
59 page-faults:u
7,817,710 cycles:u
14,067,551 instructions:u
<not supported> branches:u
56,321 branch-misses:u

```

```
0.087801579 seconds time elapsed
```

```
0.000000000 seconds user
```

```
0.064650000 seconds sys
```



- Question: Why is execution time much larger compared to PAPI measurements?

Scalasca

- ▶ Developed by FZ Jülich and TU Darmstadt
 - ▶ Some functionality became part of a framework developed with other partners
- ▶ Open source code that supports various platforms
- ▶ Support for different C, C++, and Fortran compilers
- ▶ Performance measurements require the following steps:
 1. Tools-based code instrumentation
 2. Performance measurements during code execution
 3. Analysis of performance results

Scalasca: Code Instrumentation

- ▶ Tools-based code instrumentation is meanwhile done using Score-P
 - ▶ For more information see <https://www.vi-hps.org/projects/score-p/>
- ▶ To instrument your code, prefix the compile command with `skin`
- ▶ Example:

```
% skin gcc -o ex-nopapi.x ex-nopapi.c
```


Scalasca: Code Execution

- ▶ To measure performance, prefix the execution command with `scan`

- ▶ Example:

```
% scan ./ex-nopapi.x
S=C=A=N: Scalasca 2.6 runtime summarization
S=C=A=N: ./scorep_ex-nopapi_0_sum experiment archive
S=C=A=N: Sun Jan 15 12:51:16 2023: Collect start
./ex-nopapi.x
S=C=A=N: Sun Jan 15 12:51:17 2023: Collect done (status=0) 1s
S=C=A=N: ./scorep_ex-nopapi_0_sum complete.
```

- ▶ Measurement files have been created in
`./scorep_ex-nopapi_0_sum`


- ▶ To include PAPI counter measurements, define the environment variable `SCOREP_METRIC_PAPI`:

```
% export SCOREP_METRIC_PAPI="PAPI_TOT_CYC,PAPI_LD_INS"
% skin gcc -o ex-nopapi.x ex-nopapi.c
```


Scalasca: Performance Analysis

Options for analysing statistics collected in `profile.cubex`:

- ▶ CLI `cube_stat`
- ▶ GUI `cube`



Finish with an Architecture from Turing: Bombe



[United Kingdom Government, 1945]