

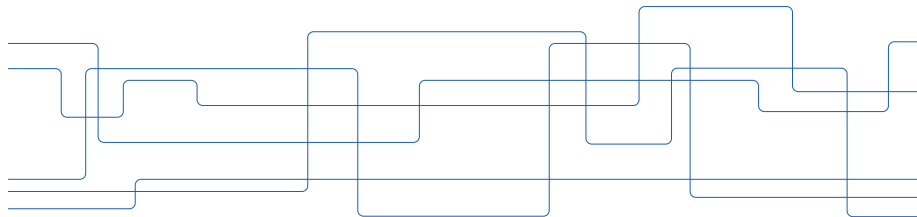


# Performance Analysis

**Dirk Pleiter**

**PDC and CST | EECS | KTH**

**January 2023**





# Overview

Introduction to Performance

Performance Measurement



# Content


Introduction to Performance

Performance Measurement

- ▶ Reliable measure of **performance** =  $\frac{\text{amount of work}}{\text{execution time}}$
- ▶ Execution time = time to execute a particular amount of work
- ▶ Ambigiouties in the definition of execution time:
  - ▶ Wall-clock time = latency to complete task
  - ▶ Resource occupation time, e.g. CPU time

# Bandwidth or Throughput vs. Latency

- ▶ Definition **bandwidth or throughput**
  - ▶ Bandwidth = amount of items passing a particular interface per time unit
  - ▶ Throughput = amount of work done in a given time
- ▶ Definition **latency (or response time)  $\Delta t$**  =  
Time between events indicating start of a task/operation and the event signaling its completion
- ▶ Example: Memory read operation




- ▶ Bandwidth: Amount of data passing memory interface per time unit
- ▶ Latency: Time from starting memory read operation until last data item arrived in register file

# Latency-Bandwidth Model

- ▶ Simple model describing **latency**  $\Delta t$  as a function of the information exchange  $I$
- ▶ Model parameters:
  - ▶ **Start-up latency**  $\lambda$  = Time between event trigger and start of arrival of response
  - ▶ **Bandwidth**  $\beta$ : Parametrizes time to exchange information  $I$
- ▶ Model:



$$\Delta t(I) = \lambda + \frac{I}{\beta}$$



# Latency-Bandwidth Model (2)

- ▶ **Effective bandwidth**  $\tilde{b} = I / \Delta t(I)$
- ▶ Example: Read  $I$  Bytes from memory or network

$$\Delta t(I) = \lambda + \frac{I}{\beta} \Rightarrow \tilde{b}(I) = \left( \frac{\lambda}{I} + \frac{1}{\beta} \right)^{-1}$$



# Information Exchange Function

- ▶ A computation implies that information is transferred from a storage device  $x$  to a storage device  $y$ .
- ▶ **Information Exchange Function:**

$I_{x,y}^k(W)$  = data transferred between computer sub-systems for specific computation  $k$

$x$  ... source storage device (e.g. memory)


$y$  ... destination storage device (e.g. register file)

$W$  ... **problem size/work-load**



# Information Exchange Model

- ▶ **Machine** = Set of interconnected devices
  - ▶ Storage devices
  - ▶ Processing/transport devices
- ▶ **Storage devices**
  - ▶ Examples: Memory, register file, cache
  - ▶ Parameters: Storage size  $\sigma$
- ▶ **Processing/transport devices**
  - ▶ Examples: Arithmetic pipeline, bus
  - ▶ Parameters: bandwidth/throughput  $\beta$ , startup latency  $\lambda$
- ▶ **Graphical representation**
  - ▶ Vertices = Storage devices
  - ▶ Edges = Processing/transport devices



# Information Exchange Model: Latency Predictions

Ansatz to predict latency:

$$\Delta t_{x,y}^k \simeq \lambda_{x,y} + l_{x,y}^k / \beta_{x,y}$$

Example:

- ▶  $x, y = R$
- ▶  $\beta_{R,R}$  = throughput arithmetic unit
- ▶  $l_{x,y}^k$  = number of input (or output) operands



**arithmetic unit**

**register file**

Beware of limitations of this ansatz:

- ▶ Transfer mechanism may depend on task size  $N$
- ▶ Bandwidth changes due to resource congestion is ignored
- ▶ ...

# Information Exchange Model: Simple Example (1/2)

- ▶ Example operation:  $y_i \leftarrow \alpha \cdot x_i$  ( $i = 1, \dots, N$ )  
 $x_i, y_i \dots$  vectors of single precision floating-point numbers  
 $\alpha \dots$  single-precision floating-point number
- ▶ Information exchange:

load $x$ , store $y$	$I_{\text{mem}} = N \cdot (4 + 4) \text{ Bytes}$
multiply $\alpha$ and $x$	$I_{\text{fp}} = N \cdot 1 \text{ Flop}$

- ▶ Assume the following hardware parameters:

Memory bandwidth	$\beta_{\text{mem}} = 1 \text{ Byte/clock cycle}$
Floating-point unit throughput	$\beta_{\text{fp}} = 1 \text{ Flop/clock cycle}$

- ▶ Latency predictions (ignoring start-up latency):

$\Delta t_{\text{mem}}$	$N \cdot 8 \text{ clock cycles}$
$\Delta t_{\text{fp}}$	$N \cdot 1 \text{ clock cycles}$

# Information Exchange Model: Simple Example (2/2)

Latency for full operation:

- ▶ No overlap of memory load/store and arithmetic operations:

$$\Delta t(N) = \Delta t_{\text{mem}} + \Delta t_{\text{fp}} = N \cdot 9 \text{ cycle}$$

- ▶ Perfect overlap of memory load/store and arithmetic operations:

$$\Delta t(N) = \max(\Delta t_{\text{mem}}, \Delta t_{\text{fp}}) = N \cdot 8 \text{ cycle}$$

👉 **Memory bandwidth limited problem**

# Arithmetic Intensity

► **Arithmetic Intensity** =

[H. Harris, 2005]

$$AI = \frac{\text{Number of floating-point operations}}{\text{Amount of transferred data}} = \frac{l_{fp}}{l_{mem}}$$

► Example:  $y_i \leftarrow \alpha \cdot x_i$

$$AI = \frac{1 \text{ Flop}}{8 \text{ Bytes}} \quad \text{single-precision}$$

$$AI = \frac{1 \text{ Flop}}{16 \text{ Bytes}} \quad \text{double-precision}$$

# Roofline Model (1/2)

[S. Williams et al., 2009]

- ▶ Floating-point and memory performance limits:

$$b_{\text{fp}} \leq B_{\text{fp}}, \quad b_{\text{mem}} \leq B_{\text{mem}}$$

- ▶ Upper limit for latency assuming perfect overlap of memory and arithmetic operations assuming the latency-bandwidth model to hold with start-up latencies:


$$\Delta t = \max \left( \frac{l_{\text{fp}}}{b_{\text{fp}}}, \frac{l_{\text{mem}}}{b_{\text{mem}}} \right) \geq \max \left( \frac{l_{\text{fp}}}{B_{\text{fp}}}, \frac{l_{\text{mem}}}{B_{\text{mem}}} \right)$$

- ▶ Upper limit (=roof) for floating-point performance (“attainable performance”)

$$b_{\text{fp}} = \frac{l_{\text{fp}}}{\Delta t} \leq \min \left( B_{\text{fp}}, \frac{l_{\text{fp}}}{l_{\text{mem}}} B_{\text{mem}} \right) = \min (B_{\text{fp}}, \text{AI} \cdot B_{\text{mem}})$$

# Roofline Model (2/2)

Example:  $y_i \leftarrow \alpha \cdot x_i$





# Content

Introduction to Performance

Performance Measurement



# Performance measurement

- ▶ Types
  - ▶ **Profiling**  
Characterize behaviour of an application in terms of aggregate performance metrics
    - ▶ Example: Average time spent in subroutine
  - ▶ **Sampling**  
Creation of profiles by means of period or random probing of performance numbers independently of the current state of application execution.
  - ▶ **Event-tracing**  
Generate list of event records comprising of time stamp, location, type information, event-specific information
- ▶ Typically additional instructions need to be inserted into program (“instrumentation”)

# Performance Measurement (2)

## ▶ Hardware support for performance measurement

- ▶ Modern processors provide hardware counters organised in a Performance Monitoring Unit (PMU)

**Hardware counters** = Small set of registers that count events

- ▶ Typical events
  - ▶ Clock tick
  - ▶ Cache miss
  - ▶ Instruction scheduled

## ▶ Issues related to performance measurement

- ▶ Code instrumentation may impact performance
- ▶ Accessibility of hardware counters
  - ▶ Typically support from kernel is required
  - ▶ System-wide services may be used for collecting data
- ▶ Management of large event traces

# Typical Time Scales

Example: Intel Nehalem architecture

Core clock cycle	300 – 500 ps
L1 cache access latency	$\sim 5$ cycles
L2 cache access latency	$\sim 10$ cycles
L3 cache access latency	$\sim 30 - 50$ cycles
Memory access latency	200 – 300 cycles

☞ Need time resolution  $\ll 1 \mu\text{sec}$

# Time Measurement

- ▶ `/usr/bin/time`
  - ▶ Wall, user, system time
  - ▶ Resolution given by ticks per second:  
`sysconf(_SC_CLK_TCK)` typically: 10 ms
- ▶ `clock_gettime(clockid_t, struct timespec *)`
  - ▶ Support for different clocks, recommended to use `CLOCK_PROCESS_CPUTIME_ID` or `CLOCK_THREAD_CPUTIME_ID`
  - ▶ Returns current clock as structure

```
struct timespec {
    time_t tv_sec;    /* seconds */
    long tv_nsec;    /* nanoseconds */
};
```
  - ▶ Include file `time.h` and link with `-lrt`

# Performance Tools: PAPI

**PAPI** = Performance Application Programming Interface

- ▶ Developed at University of Tennessee
- ▶ De facto standard to interface with PMU on various architectures
  - ▶ Examples: Intel Xeon, Xeon Phi; AMD; IBM POWER; Arm
- ▶ Selected counters supported by PAPI:

PAPI_TOT_CYC	Total cycles
PAPI_TOT_INS	Instructions completed
PAPI_FP_INS	Floating point instructions
PAPI_L1_DCM	Level 1 data cache misses
PAPI_L2_DCM	Level 2 data cache misses
PAPI_L3_DCM	Level 3 data cache misses
PAPI_BR_INS	Branch instructions
PAPI_BR_MSP	Conditional branch instructions mispredicted

# PAPI: Example Code

```
1 #include <stdlib.h>
2 #include <stdio.h>
3 #include <papi.h>
4
5 # define N 10000
6 double a[N], b[N];
7
8 void kernel() {
9     for (int i = 0; i < N; i++)
10         b[i] = a[i] + 0.1;
11 }
12
13
14 int main() {
15     // Start measurement
16     if (PAPI_hl_region_begin("kernel") != PAPI_OK)
17         exit(1);
18
19     // Some kernel
20     kernel();
21
22     // Read counters
23     if (PAPI_hl_region_end("kernel") != PAPI_OK)
24         exit(1);
25
26     return 0;
27 }
```

# PAPI: Example Execution

Compile example program:

```
% gcc -o ex.x ex.c -lpapi
```

Execute example program:

```
% export PAPI_EVENTS="PAPI_TOT_CYC,PAPI_LD_INS"  
% export PAPI_OUTPUT_DIRECTORY="."  
% ./ex.x
```

```
{  
  "papi_version": "7.0.0.0",  
  "cpu_info": "Hisilicon Kunpeng",  
  "max_cpu_rate_mhz": "2600",  
  "min_cpu_rate_mhz": "200",  
  "event_definitions": {  
    "PAPI_TOT_CYC": {  
      "component": "perf_event",  
      "type": "delta"  
    },  
    "PAPI_LD_INS": {  
      "component": "perf_event",  
      "type": "delta"  
    }  
  },  
  "threads": {  
    "0": {  
      "regions": {  
        "0": {  
          "name": "kernel",  
          "parent_region_id": "-1",  
          "cycles": "5293",  
          "real_time_nsec": "50140",  
          "PAPI_TOT_CYC": "106057",  
          "PAPI_LD_INS": "61411"  
        }  
      }  
    }  
  }  
}
```

# PAPI: Example Performance Analysis

- ▶ For  $N = 10000$  loop iterations we measured  $N_{ld} = 61411$   
 $\Rightarrow n_{ld} = N_{ld}/N \simeq 6$
- ▶ Inspection of the assembler code: Number of load instructions per loop iteration is  $n_{ld} = 6$ 
  - ▶ Tip: Use `gcc -S ex.c` to obtain the assembler code
- ▶ Inspection of the C code: Only need to load `a[i]`, i.e. the optimum would be  $n_{ld} = 1$

```

kernel:
.LFB6:
    .cfi_startproc
    sub     sp, sp, #16
    .cfi_def_cfa_offset 16
    str     wzr, [sp, 12]
    b       .L2

.L3:
    adrp    x0, a
    add     x0, x0, :lo12:a
    ldrsw   x1, [sp, 12]
    ldr     d0, [x0, x1, lsl 3]
    adrp    x0, .LC0
    ldr     d1, [x0, #:lo12:.LC0]
    fadd    d0, d0, d1
    adrp    x0, b
    add     x0, x0, :lo12:b
    ldrsw   x1, [sp, 12]
    str     d0, [x0, x1, lsl 3]
    ldr     w0, [sp, 12]
    add     w0, w0, 1
    str     w0, [sp, 12]

.L2:
    ldr     w1, [sp, 12]
    mov     w0, 9999
    cmp     w1, w0
    ble     .L3
    nop
    nop
    add     sp, sp, 16
    .cfi_def_cfa_offset 0
    ret
    
```



# PAPI: Example Execution (Optimized)

Compile example program with more optimizations enabled:

```
% gcc -O3 -o ex-03.x ex.c -lpapi
```

Execute example program:

```
% export PAPI_EVENTS="PAPI_TOT_CYC,PAPI_LD_INS"  
% export PAPI_OUTPUT_DIRECTORY="."  
% ./ex-03.x
```

```
{  
  "papi_version": "7.0.0.0",  
  "cpu_info": "Hisilicon Kunpeng",  
  "max_cpu_rate_mhz": "2600",  
  "min_cpu_rate_mhz": "200",  
  "event_definitions": {  
    "PAPI_TOT_CYC": {  
      "component": "perf_event",  
      "type": "delta"  
    },  
    "PAPI_LD_INS": {  
      "component": "perf_event",  
      "type": "delta"  
    }  
  },  
  "threads": {  
    "0": {  
      "regions": {  
        "0": {  
          "name": "kernel",  
          "parent_region_id": "-1",  
          "cycles": "2500",  
          "real_time_nsec": "22840",  
          "PAPI_TOT_CYC": "40591",  
          "PAPI_LD_INS": "6297"  
        }  
      }  
    }  
  }  
}
```

# PAPI: Example Performance Analysis (Optimized)

- ▶ For  $N = 10000$  loop iterations we measured  $N_{ld} = 6297$   
 $\Rightarrow n_{ld} = N_{ld}/N \simeq 0.5$
- ▶ Inspection of the assembler code: Number of load instructions per loop iteration is  $n_{ld} = 1$
- ▶ Question: Why do we observe less load instructions than expected?

```
kernel:
.LFB22:
    .cfi_startproc
    adrp    x0, .LC0
    adrp    x2, b
    adrp    x1, a
    mov     x3, 14464
    add     x2, x2, :lo12:b
    add     x1, x1, :lo12:a
    ldr     q1, [x0, #:lo12:.LC0]
    movk    x3, 0x1, lsl 16
    mov     x0, 0
    .p2align 3,,7
.L2:
    ldr     q0, [x1, x0]
    fadd    v0.2d, v0.2d, v1.2d
    str     q0, [x2, x0]
    add     x0, x0, 16
    cmp     x0, x3
    bne     .L2
    ret
```

# Performance Tools: Perf

**perf** = Performance analysis tool for Linux

- Uses the Performance Counters for Linux (PCL)
- Listing of available events: `perf list`
- Example usage:

```
% perf stat ./ex-03.x
```

```
Performance counter stats for './ex-03.x':
```


```

64.14 msec task-clock:u
0 context-switches:u
0 cpu-migrations:u
59 page-faults:u
7,817,710 cycles:u
14,067,551 instructions:u
<not supported> branches:u
56,321 branch-misses:u
```

```
0.087801579 seconds time elapsed
```

```
0.000000000 seconds user
```

```
0.064650000 seconds sys
```



- Question: Why is execution time much larger compared to PAPI measurements?

## Scalasca

- ▶ Developed by FZ Jülich and TU Darmstadt
  - ▶ Some functionality became part of a framework developed with other partners
- ▶ Open source code that supports various platforms
- ▶ Support for different C, C++, and Fortran compilers
- ▶ Performance measurements require the following steps:
  1. Tools-based code instrumentation
  2. Performance measurements during code execution
  3. Analysis of performance results

# Scalasca: Code Instrumentation

- ▶ Tools-based code instrumentation is meanwhile done using Score-P
  - ▶ For more information see <https://www.vi-hps.org/projects/score-p/>
- ▶ To instrument your code, prefix the compile command with `skin`
- ▶ Example:  

```
% skin gcc -o ex-nopapi.x ex-nopapi.c
```

# Scalasca: Code Execution

- ▶ To measure performance, prefix the execution command with `scan`

- ▶ Example:

```
% scan ./ex-nopapi.x
S=C=A=N: Scalasca 2.6 runtime summarization
S=C=A=N: ./scorep_ex-nopapi_0_sum experiment archive
S=C=A=N: Sun Jan 15 12:51:16 2023: Collect start
./ex-nopapi.x
S=C=A=N: Sun Jan 15 12:51:17 2023: Collect done (status=0) 1s
S=C=A=N: ./scorep_ex-nopapi_0_sum complete.
```

- ▶ Measurement files have been created in  
`./scorep_ex-nopapi_0_sum`

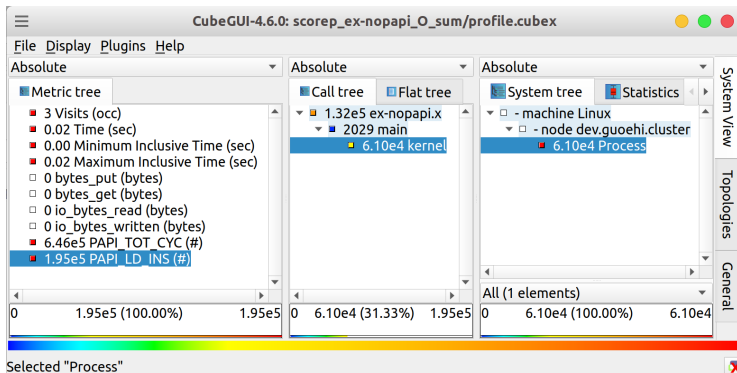
- ▶ To include PAPI counter measurements, define the environment variable `SCOREP_METRIC_PAPI`:

```
% export SCOREP_METRIC_PAPI="PAPI_TOT_CYC,PAPI_LD_INS"
% skin gcc -o ex-nopapi.x ex-nopapi.c
```


# Scalasca: Performance Analysis

Options for analysing statistics collected in `profile.cubex`:

- ▶ CLI `cube_stat`
- ▶ GUI `cube`



# Finish with an Architecture from Turing: Bombe



[United Kingdom Government, 1945]